

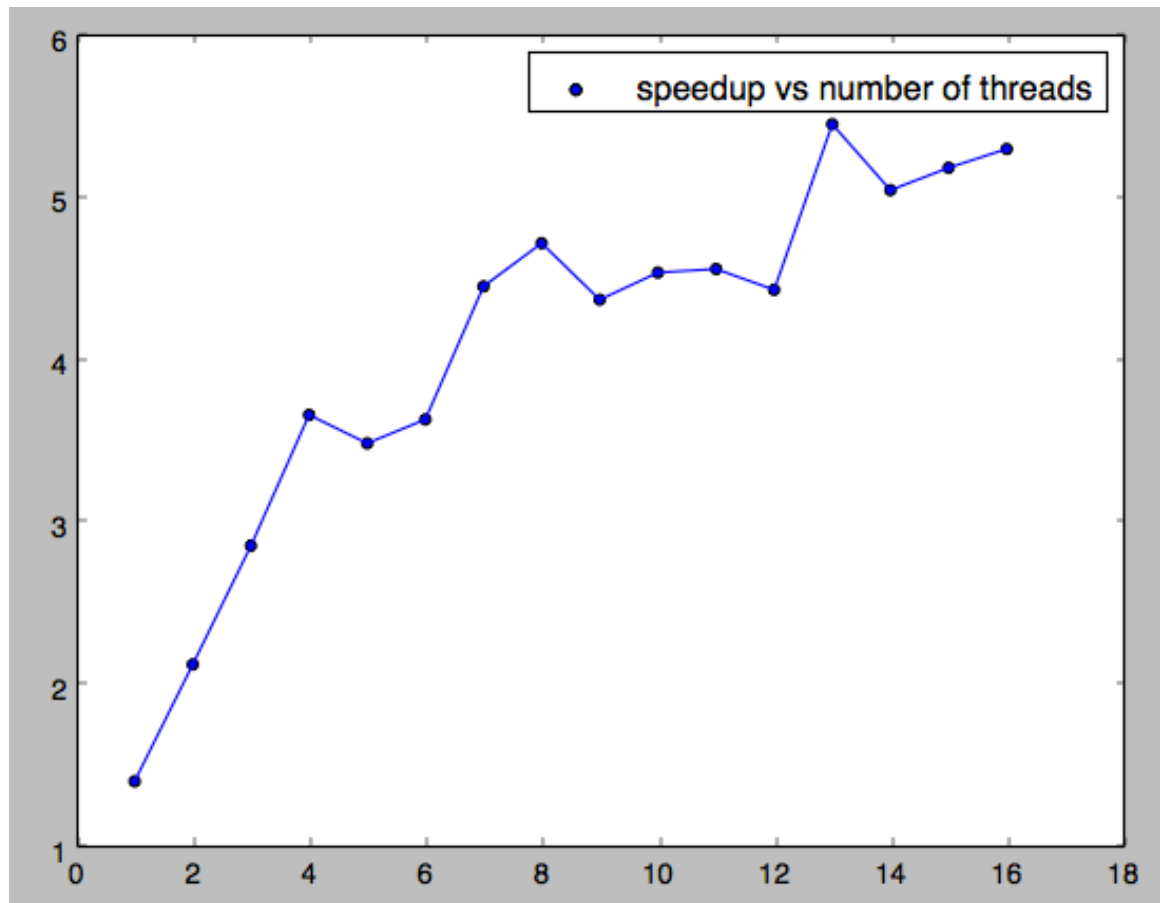
Assignment 3

2.1.1

- 1.
2. On average, my program got about a 1.8x speedup. It took a lot of the tweaking of how wide I wanted to unroll the loop, but I settled on jumping 4, rather than something wider for the radius.
3. I used `_mm_add_ps` onto my 128-bit accumulator inside the two for loops. I used `_mm_loadu_ps` to load 4 floats at a time from memory. And I used `_mm_prefetch` to look ahead. Finally, I used SSE3 instructions `_mm_hadd_ps` and `_mm_cvtss_f32` to “fold” the vector into the final “avg” value.
4. I only affected the innermost loop, since it is the bottleneck (it’s called the most times). I have one vector load per iteration, and I prefetch forward. Finally, I fold the vector and add the 4 values into a single value, and do the usual avg/num calculation.

2.3.3

1.



2.

3. I did nothing fancy. Since all of my vectorization speedup was done on the inner loop, it allowed me to just use the OpenMP `#pragma omp parallel for`

- collapse(2)” to speed up the outer loops since they’re “embarrassingly parallel.”
4. I did consider load balancing, but I figured it was not worth the overhead. It’s clear that there’s an imbalanced load, since the corners of the photo have less blurring than the direct center. A possible solution would be to first divide up the inner pixels, so that all of the jobs are CPU intensive, and equally distributed. Then a free CPU could tackle the easier jobs after it’s done.

2.3

- 1.
2. My final implementation is about 6x faster than the naïve implementation.
3. This is done by combining my OpenMP with vectorization for the inner loop blur calculation. I also did simple algorithmic enhancements, like instead of counting num for every iteration, you can just say that num has to be added for the total of (loop upper limit – loop lower limit), which saves the CPU some wasted calculation.